

Step-by-Step Guide: Building Your Own Generative AI (RAG) System

1. Define Your Use Case

- What questions should your AI answer?
- Who are your users?
- What data do you need (FAQs, manuals, reports, etc.)?

2. Collect and Prepare Data

- Gather documents (PDFs, web pages, text files, etc.)
- Clean and preprocess text (remove noise, fix encoding)
- Organize data into a consistent format (markdown, JSON, etc.)

3. Build a Knowledge Base

- Split documents into chunks (e.g., paragraphs, sections)
- Store each chunk with metadata (title, source, etc.)
- Save as markdown or JSON for easy processing

4. Generate Embeddings & Create a Vector Store

- Choose an embedding model (e.g., `sentence-transformers/all-MiniLM-L6-v2`)
- Convert each chunk to a vector (embedding)
- Store vectors in a vector database (e.g., FAISS, Chroma, Pinecone)

5. Integrate a Language Model (LLM)

- Choose your LLM: OpenAI (cloud), Ollama (local), HuggingFace, etc.
- Set up the LLM to answer questions using retrieved context
- Use frameworks like LangChain for easy integration

6. Build a Retrieval-Augmented Generation (RAG) Pipeline

- On each user question:
 1. Embed the question
 2. Retrieve top-k relevant chunks from the vector store
 3. Pass context + question to the LLM for answer generation

7. Expose Your System via API or Web Interface

- Use FastAPI or Flask to create a REST API
- Optionally, build a web chat interface (React, Streamlit, etc.)
- Log all interactions for monitoring and improvement

8. Add MLOps: Experiment Tracking & Data Versioning

- Use MLflow to track:
 - Latency, token usage, retrieval quality, etc.
 - Model versions and parameters
- Use DVC to version your data and knowledge base

9. Test and Evaluate

- Create a set of test questions
- Measure accuracy, latency, and user satisfaction
- Iterate on data, retrieval, and prompt engineering

10. Deploy and Monitor

- Deploy API/web app to cloud or on-premises
- Monitor health, usage, and errors
- Regularly update data and retrain embeddings as needed

Example Tech Stack

- **Data Processing:** Python, BeautifulSoup, Pandas
 - **Embeddings:** sentence-transformers, HuggingFace
 - **Vector Store:** FAISS, Chroma, Pinecone
 - **LLM:** Ollama (Llama 3), OpenAI GPT, HuggingFace Transformers
 - **API:** FastAPI, Flask
 - **Web UI:** React, Streamlit, Flask templates
 - **MLOps:** MLflow, DVC, Git
 - **Deployment:** Docker, cloud VM, on-prem server
-

Sample Project Structure

```
project-root/
|-- data/
|   |-- raw/           # Source documents
|   `-- knowledge_base/ # Processed chunks/snippets
|-- scripts/           # All utility scripts
|-- app/               # API and web app code
|-- models/            # LLMs, embeddings
|-- mlflow.db           # MLflow tracking
|-- .dvc/               # DVC config
|-- requirements.txt
`-- README.md
```

Tips for Success

- Start small: prototype with a few documents and questions
 - Use open-source models for privacy and cost control
 - Log everything: questions, answers, retrievals, errors
 - Regularly update your knowledge base and embeddings
 - Involve users for feedback and improvement
-

Resources

- [LangChain Documentation](#)
 - [MLflow Documentation](#)
 - [DVC Documentation](#)
 - [Ollama \(local LLMs\)](#)
 - [HuggingFace Transformers](#)
 - [FastAPI](#)
-

Prepared by your AI coding assistant